

Unit 4 Challenge - Tier 3

November 6, 2025

1 Springboard Data Science Career Track Unit 4 Challenge - Tier 3 Complete

1.1 Objectives

Hey! Great job getting through those challenging DataCamp courses. You're learning a lot in a short span of time.

In this notebook, you're going to apply the skills you've been learning, bridging the gap between the controlled environment of DataCamp and the *slightly* messier work that data scientists do with actual datasets!

Here's the mystery we're going to solve: *which boroughs of London have seen the greatest increase in housing prices, on average, over the last two decades?*

A borough is just a fancy word for district. You may be familiar with the five boroughs of New York... well, there are 32 boroughs within Greater London ([here's some info for the curious](#)). Some of them are more desirable areas to live in, and the data will reflect that with a greater rise in housing prices.

This is the Tier 3 notebook, which means it's not filled in at all: we'll just give you the skeleton of a project, the brief and the data. It's up to you to play around with it and see what you can find out! Good luck! If you struggle, feel free to look at easier tiers for help; but try to dip in and out of them, as the more independent work you do, the better it is for your learning!

This challenge will make use of only what you learned in the following DataCamp courses:

- Prework courses (Introduction to Python for Data Science, Intermediate Python for Data Science)
- Data Types for Data Science - Python Data Science Toolbox (Part One) - pandas Foundations
- Manipulating DataFrames with pandas - Merging DataFrames with pandas

Of the tools, techniques and concepts in the above DataCamp courses, this challenge should require the application of the following:

- **pandas** - **data ingestion and inspection** (pandas Foundations, Module One)
- **exploratory data analysis** (pandas Foundations, Module Two)
- **tidying and cleaning** (Manipulating DataFrames with pandas, Module Three)
- **transforming DataFrames** (Manipulating DataFrames with pandas, Module One)
- **subsetting DataFrames with lists** (Manipulating DataFrames with pandas, Module One)
- **filtering DataFrames** (Manipulating DataFrames with pandas, Module One)
- **grouping data** (Manipulating DataFrames with pandas, Module Four)
- **melting data** (Manipulating DataFrames with pandas, Module Three)
- **advanced indexing** (Manipulating DataFrames with pandas, Module Four)
- **matplotlib** (Intermediate Python for Data Science, Module One)
- **fundamental data types** (Data Types for

Data Science, Module One) - **dictionaries** (Intermediate Python for Data Science, Module Two) - **handling dates and times** (Data Types for Data Science, Module Four) - **function definition** (Python Data Science Toolbox - Part One, Module One) - **default arguments, variable length, and scope** (Python Data Science Toolbox - Part One, Module Two) - **lambda functions and error handling** (Python Data Science Toolbox - Part One, Module Four)

1.2 The Data Science Pipeline

This is Tier Three, so we'll get you started. But after that, it's all in your hands! When you feel done with your investigations, look back over what you've accomplished, and prepare a quick presentation of your findings for the next mentor meeting.

Data Science is magical. In this case study, you'll get to apply some complex machine learning algorithms. But as [David Spiegelhalter](#) reminds us, there is no substitute for simply **taking a really, really good look at the data**. Sometimes, this is all we need to answer our question.

Data Science projects generally adhere to the four stages of Data Science Pipeline: 1. Sourcing and loading 2. Cleaning, transforming, and visualizing 3. Modeling 4. Evaluating and concluding

1.2.1 1. Sourcing and Loading

Any Data Science project kicks off by importing **pandas**. The documentation of this wonderful library can be found [here](#). As you've seen, pandas is conveniently connected to the [Numpy](#) and [Matplotlib](#) libraries.

Hint: This part of the data science pipeline will test those skills you acquired in the pandas Foundations course, Module One.

1.1. Importing Libraries

```
[312]: # Let's import the pandas, numpy libraries as pd, and np respectively.  
import numpy as np  
import pandas as pd  
  
# Load the pyplot collection of functions from matplotlib, as plt  
import matplotlib.pyplot as plt
```

1.2. Loading the data Your data comes from the [London Datastore](#): a free, open-source data-sharing portal for London-oriented datasets.

```
[313]: # First, make a variable called url_LondonHousePrices, and assign it the  
# following link, enclosed in quotation-marks as a string:  
# https://data.london.gov.uk/download/uk-house-price-index/  
# 70ac0766-8902-4eb5-aab5-01951aaed773/UK%20House%20price%20index.xls  
  
url_LondonHousePrices = "https://data.london.gov.uk/download/  
uk-house-price-index/70ac0766-8902-4eb5-aab5-01951aaed773/  
UK%20House%20price%20index.xls"
```

```

# The dataset we're interested in contains the Average prices of the houses, and is actually on a particular sheet of the Excel file.
# As a result, we need to specify the sheet name in the read_excel() method.
# Put this data into a variable called properties.
properties = pd.read_excel(url_LondonHousePrices, sheet_name='Average price', index_col= None)

```

1.2.2 2. Cleaning, transforming, and visualizing

This second stage is arguably the most important part of any Data Science project. The first thing to do is take a proper look at the data. Cleaning forms the majority of this stage, and can be done both before or after Transformation.

The end goal of data cleaning is to have tidy data. When data is tidy:

1. Each variable has a column.
2. Each observation forms a row.

Keep the end goal in mind as you move through this process, every step will take you closer.

Hint: This part of the data science pipeline should test those skills you acquired in: - Intermediate Python for data science, all modules. - pandas Foundations, all modules. - Manipulating DataFrames with pandas, all modules. - Data Types for Data Science, Module Four. - Python Data Science Toolbox - Part One, all modules

2.1. Exploring your data

Think about your pandas functions for checking out a dataframe.

```
[314]: print(properties.info())
print(properties.head())
#print(properties.tail())
print(properties.shape)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 369 entries, 0 to 368
Data columns (total 49 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Unnamed: 0        368 non-null    datetime64[ns]
 1   City of London    369 non-null    object  
 2   Barking & Dagenham 369 non-null    object  
 3   Barnet            369 non-null    object  
 4   Bexley            369 non-null    object  
 5   Brent             369 non-null    object  
 6   Bromley           369 non-null    object  
 7   Camden            369 non-null    object  
 8   Croydon           369 non-null    object  
 9   Ealing             369 non-null    object  
 10  Enfield           369 non-null    object  
 11  Greenwich          369 non-null    object 

```

```

12 Hackney          369 non-null   object
13 Hammersmith & Fulham 369 non-null   object
14 Haringey         369 non-null   object
15 Harrow           369 non-null   object
16 Havering         369 non-null   object
17 Hillingdon       369 non-null   object
18 Hounslow          369 non-null   object
19 Islington        369 non-null   object
20 Kensington & Chelsea 369 non-null   object
21 Kingston upon Thames 369 non-null   object
22 Lambeth          369 non-null   object
23 Lewisham          369 non-null   object
24 Merton            369 non-null   object
25 Newham            369 non-null   object
26 Redbridge         369 non-null   object
27 Richmond upon Thames 369 non-null   object
28 Southwark         369 non-null   object
29 Sutton            369 non-null   object
30 Tower Hamlets    369 non-null   object
31 Waltham Forest   369 non-null   object
32 Wandsworth        369 non-null   object
33 Westminster      369 non-null   object
34 Unnamed: 34        0 non-null     float64
35 Inner London      369 non-null   object
36 Outer London      369 non-null   object
37 Unnamed: 37        0 non-null     float64
38 NORTH EAST        369 non-null   object
39 NORTH WEST        369 non-null   object
40 YORKS & THE HUMBER 369 non-null   object
41 EAST MIDLANDS    369 non-null   object
42 WEST MIDLANDS    369 non-null   object
43 EAST OF ENGLAND   369 non-null   object
44 LONDON            369 non-null   object
45 SOUTH EAST        369 non-null   object
46 SOUTH WEST        369 non-null   object
47 Unnamed: 47        0 non-null     float64
48 England           369 non-null   object

```

dtypes: datetime64[ns](1), float64(3), object(45)

memory usage: 141.4+ KB

None

	Unnamed: 0	City of London	Barking & Dagenham	Barnet	Bexley	\
0	NaT	E09000001	E09000002	E09000003	E09000004	
1	1995-01-01	90347	51870	98948	64956	
2	1995-02-01	81213	52513	98848	64786	
3	1995-03-01	78168	52701	97848	64366	
4	1995-04-01	76172	54618	96273	64276	

Brent Bromley Camden Croydon Ealing ... NORTH WEST \

```

0   E09000005   E09000006   E09000007   E09000008   E09000009   ...   E12000002
1       76880      83082      119775      70118      85469   ...   40907
2       77651      83068      118365      69908      86551   ...   40877
3       77644      82856      119131      69666      87067   ...   41351
4       78668      82525      118948      69562      87933   ...   41195

    YORKS & THE HUMBER EAST MIDLANDS WEST MIDLANDS EAST OF ENGLAND      LONDON \
0           E12000003      E12000004      E12000005      E12000006      E12000007
1           42171        43856        46470        56098        79687
2           41912        44344        47249        55991        77913
3           42544        43701        47345        55574        79110
4           42934        44414        47359        55966        79708

    SOUTH EAST SOUTH WEST Unnamed: 47      England
0   E12000008   E12000009          NaN  E92000001
1       64502      52799          NaN      50231
2       64196      52462          NaN      50130
3       64597      51716          NaN      50229
4       65111      52877          NaN      50597

[5 rows x 49 columns]
(369, 49)

```

2.2. Cleaning the data

You might find you need to transpose your dataframe, check out what its row indexes are, and reset the index. You also might find you need to assign the values of the first row to your column headings . (Hint: recall the .columns feature of DataFrames, as well as the iloc[] method).

Don't be afraid to use StackOverflow for help with this.

2.3. Cleaning the data (part 2)

You might we have to **rename** a couple columns. How do you do this? The clue's pretty bold...

```
[393]: properties_transposed = properties.T.reset_index()
print(properties_transposed.head(3))
```

```

      index          0          1          2 \
0      Unnamed: 0      NaT  1995-01-01 00:00:00  1995-02-01 00:00:00
1  City of London  E09000001                  90347          81213
2  Barking & Dagenham  E09000002                  51870          52513

      3          4          5 \
0  1995-03-01 00:00:00  1995-04-01 00:00:00  1995-05-01 00:00:00
1          78168          76172          83392
2          52701          54618          54524

      6          7          8 ... \
0  1995-06-01 00:00:00  1995-07-01 00:00:00  1995-08-01 00:00:00 ...
1          93757          108801          110976 ...
```

2	55200	53569	53691	...
	359	360	361	\
0	2024-11-01 00:00:00	2024-12-01 00:00:00	2025-01-01 00:00:00	
1	728756	770014	789031	
2	350798	354406	356394	
	362	363	364	\
0	2025-02-01 00:00:00	2025-03-01 00:00:00	2025-04-01 00:00:00	
1	749626	715511	747919	
2	361644	369633	371161	
	365	366	367	\
0	2025-05-01 00:00:00	2025-06-01 00:00:00	2025-07-01 00:00:00	
1	802389	807327	780192	
2	367979	362254	364858	
	368			
0	2025-08-01 00:00:00			
1	711729			
2	362062			

[3 rows x 370 columns]

```
[316]: properties_col = properties_transposed.iloc[0]
new_prop = pd.DataFrame(properties_transposed.values[1:], columns = ↴
    ↪properties_col)
print(new_prop.head())
new_prop.columns
```

0	Unnamed: 0	NaT	1995-01-01 00:00:00	1995-02-01 00:00:00	\
0	City of London	E09000001	90347	81213	
1	Barking & Dagenham	E09000002	51870	52513	
2	Barnet	E09000003	98948	98848	
3	Bexley	E09000004	64956	64786	
4	Brent	E09000005	76880	77651	
0	1995-03-01 00:00:00	1995-04-01 00:00:00	1995-05-01 00:00:00	\	
0	78168	76172	83392		
1	52701	54618	54524		
2	97848	96273	95737		
3	64366	64276	63995		
4	77644	78668	79464		
0	1995-06-01 00:00:00	1995-07-01 00:00:00	1995-08-01 00:00:00	...	\
0	93757	108801	110976	...	
1	55200	53569	53691	...	
2	95577	96992	97968	...	

```

3          64251          63721          64431 ...
4          80118          79920          79295 ...

0 2024-11-01 00:00:00 2024-12-01 00:00:00 2025-01-01 00:00:00 \
0          728756          770014          789031
1          350798          354406          356394
2          628891          621602          611125
3          405701          403968          403383
4          590969          579308          567742

0 2025-02-01 00:00:00 2025-03-01 00:00:00 2025-04-01 00:00:00 \
0          749626          715511          747919
1          361644          369633          371161
2          599845          597110          591360
3          405699          405705          398795
4          564953          564650          567344

0 2025-05-01 00:00:00 2025-06-01 00:00:00 2025-07-01 00:00:00 \
0          802389          807327          780192
1          367979          362254          364858
2          583161          592221          601498
3          400821          402247          411080
4          557669          543691          536131

0 2025-08-01 00:00:00
0          711729
1          362062
2          615503
3          412151
4          560535

```

[5 rows x 370 columns]

```
[316]: Index(['Unnamed: 0', NaT, 1995-01-01 00:00:00,
   1995-02-01 00:00:00, 1995-03-01 00:00:00, 1995-04-01 00:00:00,
   1995-05-01 00:00:00, 1995-06-01 00:00:00, 1995-07-01 00:00:00,
   1995-08-01 00:00:00,
   ...
   2024-11-01 00:00:00, 2024-12-01 00:00:00, 2025-01-01 00:00:00,
   2025-02-01 00:00:00, 2025-03-01 00:00:00, 2025-04-01 00:00:00,
   2025-05-01 00:00:00, 2025-06-01 00:00:00, 2025-07-01 00:00:00,
   2025-08-01 00:00:00],
  dtype='object', name=0, length=370)
```

2.4.Transforming the data

Remember what Wes McKinney said about tidy data?

You might need to **melt** your DataFrame here.

```
[324]: new_prop.rename(columns = {new_prop.columns[0]: 'London_boroughs', new_prop.
    ↪columns[1]: 'ID'}, inplace = True)
properties_melt = pd.melt(new_prop, id_vars = ['London_boroughs', 'ID'], ↪
    ↪var_name = 'Date', value_name = 'Avg_price')
print(properties_melt.head())
```

	London_boroughs	ID	Date	Avg_price
0	City of London	E09000001	1995-01-01 00:00:00	90347
1	Barking & Dagenham	E09000002	1995-01-01 00:00:00	51870
2	Barnet	E09000003	1995-01-01 00:00:00	98948
3	Bexley	E09000004	1995-01-01 00:00:00	64956
4	Brent	E09000005	1995-01-01 00:00:00	76880

Remember to make sure your column data types are all correct. Average prices, for example, should be floating point numbers...

```
[325]: properties_melt['Date'] = pd.to_datetime(properties_melt['Date'])
properties_melt['Avg_price'] = properties_melt['Avg_price'].astype(float)
print(properties_melt.info())
print(properties_melt.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17664 entries, 0 to 17663
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   London_boroughs  17664 non-null   object 
 1   ID                16560 non-null   object 
 2   Date              17664 non-null   datetime64[ns]
 3   Avg_price         16560 non-null   float64
dtypes: datetime64[ns](1), float64(1), object(2)
memory usage: 552.1+ KB
None
      London_boroughs     ID       Date  Avg_price
0   City of London  E09000001 1995-01-01    90347.0
1   Barking & Dagenham E09000002 1995-01-01    51870.0
2   Barnet          E09000003 1995-01-01    98948.0
3   Bexley          E09000004 1995-01-01    64956.0
4   Brent           E09000005 1995-01-01    76880.0
```

2.5. Cleaning the data (part 3)

Do we have an equal number of observations in the ID, Average Price, Month, and London Borough columns? Remember that there are only 32 London Boroughs. How many entries do you have in that column?

Check out the contents of the London Borough column, and if you find null values, get rid of them however you see fit.

```
[332]: print(properties_melt.isna().any())
print(properties_melt.isna().sum())
print(properties_melt.shape)
prop_clean = properties_melt.dropna()
print(prop_clean.isna().any())
```

```
London_boroughs      False
ID                  True
Date                False
Avg_price          True
dtype: bool
London_boroughs      0
ID                  1104
Date                0
Avg_price          1104
dtype: int64
(17664, 4)
London_boroughs      False
ID                  False
Date                False
Avg_price          False
dtype: bool
```

2.6. Visualizing the data

To visualize the data, why not subset on a particular London Borough? Maybe do a line plot of Month against Average Price?

```
[333]: prop_clean.loc[:, 'Year'] = prop_clean['Date'].dt.year
prop_clean.loc[:, 'Month'] = prop_clean['Date'].dt.month
print(prop_clean.head())
```

	London_boroughs	ID	Date	Avg_price	Year	Month
0	City of London	E09000001	1995-01-01	90347.0	1995	1
1	Barking & Dagenham	E09000002	1995-01-01	51870.0	1995	1
2	Barnet	E09000003	1995-01-01	98948.0	1995	1
3	Bexley	E09000004	1995-01-01	64956.0	1995	1
4	Brent	E09000005	1995-01-01	76880.0	1995	1

```
/var/folders/hq/kk6s209d61dcscgjxy9bp4tw0000gn/T/ipykernel_57889/997687871.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
prop_clean.loc[:, 'Year'] = prop_clean['Date'].dt.year
/var/folders/hq/kk6s209d61dcscgjxy9bp4tw0000gn/T/ipykernel_57889/997687871.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
prop_clean.loc[:, 'Month'] = prop_clean['Date'].dt.month
```

To limit the number of data points you have, you might want to extract the year from every month value your *Month* column.

To this end, you *could* apply a *lambda function*. Your logic could work as follows: 1. look through the *Month* column 2. extract the year from each individual value in that column 3. store that corresponding year as separate column.

Whether you go ahead with this is up to you. Just so long as you answer our initial brief: which boroughs of London have seen the greatest house price increase, on average, over the past two decades?

```
[380]: from scipy import stats
desc_stat = prop_clean.groupby(['London_boroughs', 'Year'])['Avg_price'].
    agg(mean_price = 'mean', std_error =
        lambda x: stats.sem(x, ddof = 1)).reset_index()
print(desc_stat)
```

	London_boroughs	Year	mean_price	std_error
0	Barking & Dagenham	1995	53265.416667	305.060040
1	Barking & Dagenham	1996	53162.666667	266.963462
2	Barking & Dagenham	1997	57537.583333	338.416017
3	Barking & Dagenham	1998	61969.666667	577.836250
4	Barking & Dagenham	1999	67145.333333	467.053552
...
1390	YORKS & THE HUMBER	2021	175033.833333	1590.090939
1391	YORKS & THE HUMBER	2022	191155.833333	1810.666497
1392	YORKS & THE HUMBER	2023	192742.500000	835.833140
1393	YORKS & THE HUMBER	2024	196979.000000	1291.157057
1394	YORKS & THE HUMBER	2025	203766.375000	1369.595377

```
[1395 rows x 4 columns]
```

```
[ ]:
```

```
[ ]:
```

```
[384]: bark_dag = desc_stat[(desc_stat['London_boroughs'] == 'Barking & Dagenham') &_
    ~(desc_stat['Year'].between(2000, 2020))]
print(bark_dag.head())
hamm_ful = desc_stat[(desc_stat['London_boroughs'] == 'Hammersmith & Fulham') &_
    ~(desc_stat['Year'].between(2000, 2020))]
print(hamm_ful.head())
```

```

enfield = desc_stat[(desc_stat['London_boroughs'] == 'Enfield') &
    ↪(desc_stat['Year'].between(2000, 2020))]
print(enfield.head())
sutt = desc_stat[(desc_stat['London_boroughs'] == 'Sutton') &
    ↪(desc_stat['Year'].between(2000, 2020))]
print(sutt.head())

```

	London_boroughs	Year	mean_price	std_error
5	Barking & Dagenham	2000	79715.500000	1244.078109
6	Barking & Dagenham	2001	91140.250000	1521.845516
7	Barking & Dagenham	2002	115356.250000	3267.367108
8	Barking & Dagenham	2003	146478.750000	1108.238283
9	Barking & Dagenham	2004	162593.666667	1553.686483
	London_boroughs	Year	mean_price	std_error
470	Hammersmith & Fulham	2000	256096.333333	3490.371241
471	Hammersmith & Fulham	2001	281629.416667	2424.467528
472	Hammersmith & Fulham	2002	313260.333333	5366.348799
473	Hammersmith & Fulham	2003	330190.500000	2635.816530
474	Hammersmith & Fulham	2004	355643.500000	3361.576654
	London_boroughs	Year	mean_price	std_error
346	Enfield	2000	129405.916667	1769.780659
347	Enfield	2001	144141.333333	2112.194415
348	Enfield	2002	169714.500000	4034.296254
349	Enfield	2003	202408.000000	1321.375781
350	Enfield	2004	218631.750000	2097.893298
	London_boroughs	Year	mean_price	std_error
1183	Sutton	2000	127332.250000	1991.120117
1184	Sutton	2001	139421.583333	1943.514122
1185	Sutton	2002	164307.750000	3150.456776
1186	Sutton	2003	189874.083333	1031.409322
1187	Sutton	2004	204041.083333	2415.975073

```

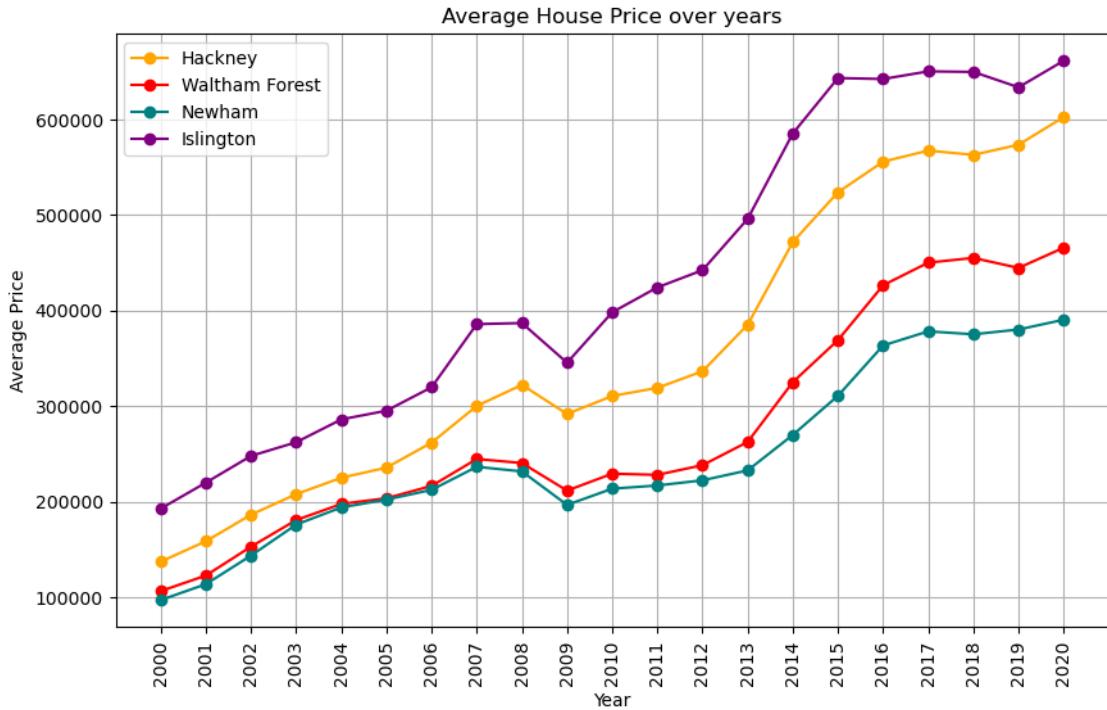
[391]: hackney = desc_stat[(desc_stat['London_boroughs'] == 'Hackney') &
    ↪(desc_stat['Year'].between(2000, 2020))]
print(hackney.head())
walth = desc_stat[(desc_stat['London_boroughs'] == 'Waltham Forest') &
    ↪(desc_stat['Year'].between(2000, 2020))]
print(walth.head())
new = desc_stat[(desc_stat['London_boroughs'] == 'Newham') & (desc_stat['Year'].
    ↪between(2000, 2020))]
print(new.head())
isl = desc_stat[(desc_stat['London_boroughs'] == 'Islington') &
    ↪(desc_stat['Year'].between(2000, 2020))]
print(isl.head())

ax = hackney.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', ↪
    ↪color = "orange", figsize = (10,6), label = "Hackney")

```

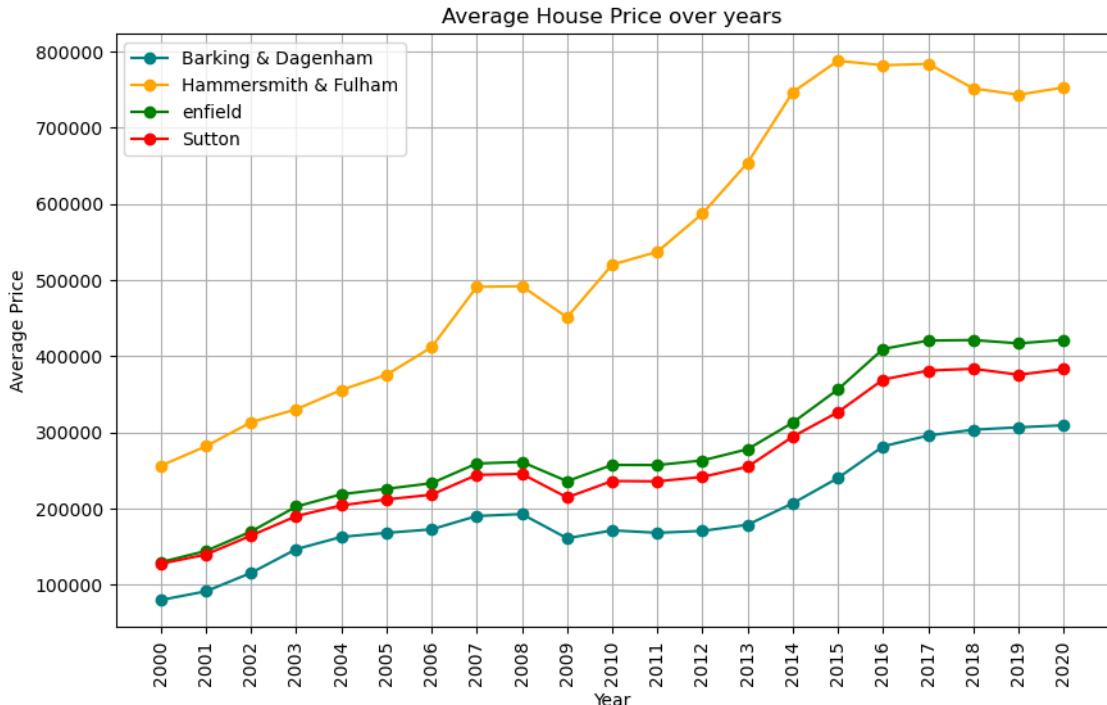
```
walth.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "red",ax = ax ,figsize = (10,6), label = "Waltham Forest")
new.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "teal",ax = ax ,figsize = (10,6), label = "Newham")
isl.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "purple",ax = ax ,figsize = (10,6), label = "Islington")
plt.xlabel('Year')
plt.ylabel('Average Price')
plt.title(f'Average House Price over years')
plt.xticks(range(2000, 2021, 1), rotation = 90)
plt.grid(True)
```

	London_boroughs	Year	mean_price	std_error
439	Hackney	2000	137228.583333	2633.154032
440	Hackney	2001	158608.583333	2244.104655
441	Hackney	2002	186382.916667	4015.592651
442	Hackney	2003	207890.250000	752.405185
443	Hackney	2004	225058.833333	2504.723035
	London_boroughs	Year	mean_price	std_error
1276	Waltham Forest	2000	106384.166667	1675.302772
1277	Waltham Forest	2001	122711.583333	2701.200725
1278	Waltham Forest	2002	152959.083333	3813.146113
1279	Waltham Forest	2003	180774.250000	876.902876
1280	Waltham Forest	2004	197606.750000	1512.158377
	London_boroughs	Year	mean_price	std_error
966	Newham	2000	97050.916667	1672.353914
967	Newham	2001	113638.250000	2076.717231
968	Newham	2002	143616.916667	3772.803241
969	Newham	2003	175987.416667	1590.765184
970	Newham	2004	194003.416667	1923.681924
	London_boroughs	Year	mean_price	std_error
687	Islington	2000	192760.000000	2986.830508
688	Islington	2001	219962.166667	2945.134866
689	Islington	2002	247923.333333	3196.264143
690	Islington	2003	262256.750000	1902.011727
691	Islington	2004	286134.250000	2145.673802



```
[383]: ax = bark_dag.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "teal", figsize = (10,6), label = "Barking & Dagenham")
hamm_ful.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "orange",ax = ax ,figsize = (10,6), label = "Hammersmith & Fulham")
enfield.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "green",ax = ax ,figsize = (10,6), label = "enfield")
sutt.plot(x = 'Year', y = 'mean_price' , kind = "line", marker = 'o', color = "red",ax = ax ,figsize = (10,6), label = "Sutton")
plt.xlabel('Year')
plt.ylabel('Average Price')
plt.title(f'Average House Price over years')
plt.xticks(range(2000, 2021, 1), rotation = 90)
plt.grid(True)

plt.show()
```



3. Modeling

Consider creating a function that will calculate a ratio of house prices, comparing the price of a house in 2018 to the price in 1998.

Consider calling this function create_price_ratio.

You'd want this function to:

1. Take a filter of dfg, specifically where this filter constrains the London_Borough, as an argument. For example, one admissible argument should be: `dfg[dfg['London_Borough']=='Camden']`.
2. Get the Average Price for that Borough, for the years 1998 and 2018.
3. Calculate the ratio of the Average Price for 1998 divided by the Average Price for 2018.
4. Return that ratio.

Once you've written this function, you ultimately want to use it to iterate through all the unique London_Boroughs and work out the ratio capturing the difference of house prices between 1998 and 2018.

Bear in mind: you don't have to write a function like this if you don't want to. If you can solve the brief otherwise, then great!

Hint: This section should test the skills you acquired in: - Python Data Science Toolbox - Part One, all modules

```
[291]: def cal_price_ratio(desc_stat, London_boroughs, start_year, end_year):
    ratios = []
    for borough, group in desc_stat.groupby('London_boroughs'):
        start_price = group.loc[group['Year'] == start_year, 'Avg_price'].values[0]
        end_price = group.loc[group['Year'] == end_year, 'Avg_price'].values[0]
        ratio = start_price / end_price
        ratios.append(ratio)
    return ratios
```

```

    end_price = group.loc[group['Year'] == end_year, 'Avg_price'].values
    price_ratio = start_price[0]/end_price[0]
    ratios.append((borough, price_ratio))
return pd.DataFrame(ratios, columns = ['London_boroughs', 'price_ratio'])

```

[293]:

```

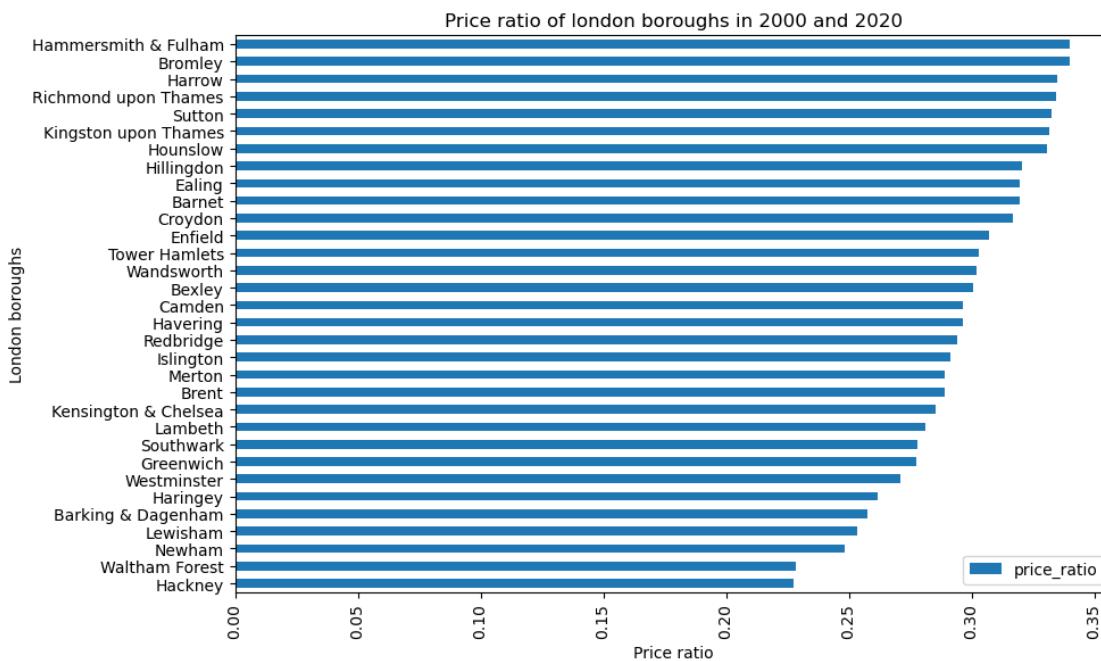
London_boroughs = desc_stat['London_boroughs']
ratio_df = cal_price_ratio(desc_stat, London_boroughs, start_year = 2000, end_year = 2020)
#print(ratio_df)
filter_ratio_df = ratio_df.query("London_boroughs not in [ 'City of London', 'SOUTH EAST', 'SOUTH WEST', 'WEST MIDLANDS', 'YORKS & THE HUMBER', 'EAST MIDLANDS', 'EAST OF ENGLAND', 'NORTH EAST', 'NORTH WEST', 'LONDON', 'Inner London', 'Outer London', 'England']")
filter_df = filter_ratio_df.sort_values(by = 'price_ratio', ascending = True)
print(filter_df.reset_index(drop = True))

```

	London_boroughs	price_ratio
0	Hackney	0.227705
1	Waltham Forest	0.228332
2	Newham	0.248508
3	Lewisham	0.253564
4	Barking & Dagenham	0.257780
5	Haringey	0.261997
6	Westminster	0.270863
7	Greenwich	0.277400
8	Southwark	0.277964
9	Lambeth	0.281007
10	Kensington & Chelsea	0.285334
11	Brent	0.289047
12	Merton	0.289254
13	Islington	0.291334
14	Redbridge	0.294050
15	Havering	0.296365
16	Camden	0.296453
17	Bexley	0.300549
18	Wandsworth	0.302018
19	Tower Hamlets	0.303156
20	Enfield	0.307079
21	Croydon	0.317038
22	Barnet	0.319449
23	Ealing	0.319566
24	Hillingdon	0.320650
25	Hounslow	0.330958
26	Kingston upon Thames	0.331777
27	Sutton	0.332547
28	Richmond upon Thames	0.334519
29	Harrow	0.334861

30	Bromley	0.339794
31	Hammersmith & Fulham	0.340128

```
[306]: filter_df.plot(x = "London_boroughs" , y = "price_ratio" , kind = "barh",  
                     figsize = (10,6))  
plt.title("Price ratio of london boroughs in 2000 and 2020")  
plt.xlabel("Price ratio")  
plt.ylabel("London boroughs")  
plt.xticks(rotation=90)  
plt.tight_layout()  
plt.show()
```



1.2.3 4. Conclusion

What can you conclude? Type out your conclusion below.

Look back at your notebook. Think about how you might summarize what you have done, and prepare a quick presentation on it to your mentor at your next meeting.

We hope you enjoyed this practical project. It should have consolidated your data hygiene and pandas skills by looking at a real-world problem involving just the kind of dataset you might encounter as a budding data scientist. Congratulations, and looking forward to seeing you at the next step in the course!

```
[ ]: """Based on the price ratio calculated between the years 2000 and 2020, Hackney  
and Waltham Forest have the lowest ratios,
```

*0.2277 and 0.2283 respectively. Since a smaller price ratio indicates a greater
↳ increase in average house price, these two boroughs
have experienced the largest growth in house prices over the past two decades. □
↳ " " "*